

Q2 — Program Charter

This is the complete written response for **Q2**. The presentation uses a concise summary and links back to this evidence.

Submission baseline: v0.3.8 · execution 11901 · RUN-S2-11902-16e7090e

Supplementary demonstration evidence: execution 11958 · RUN-S2-11959-16e7090e

Published view: <https://vjra.us/prd-genie.html#q2-program-charter>

Establish a controlled path from validated candidate to measurable adoption

Enable product teams to transform fragmented stakeholder evidence into consistent, traceable, and planning-ready documentation. PRD Orchestrator reduces administrative effort while preserving source evidence, ambiguity, stakeholder decisions, and human accountability. It supports PM and TPM judgment rather than replacing it.

Charter identity

- Organization: NeuronForge Technologies
- Program owner: AI PM/TPM · Product & Innovation
- Charter status: Draft for pilot approval
- Candidate baseline: v0.3.8
- Operating model: Manual trigger · governed release
- Executive sponsor: TBD · confirmation required

Scope

- Sources: transcripts, briefs, notes, clarifications, and templates.
- Intelligence: extraction, gap analysis, PRD, and story breakdown.
- Governance: human approval, Agreement Gate, and fail-closed export.
- Planning: feature/story criteria, priorities, and advisory sizing.
- Assurance: traceability, code controls, and LLM evaluation.
- Operations: Google Drive delivery, evidence, tokens, and cost.

Non-scope

- Autonomous contradiction resolution or approval.
- Replacing engineering estimates or commitments.
- Automatic Jira creation or production backlog updates.
- Company-wide deployment before pilot evidence.
- Fine-tuning before workflow optimization.
- Scheduled production triggering during the pilot.

Success measures

Measure	Target	Current position	Classification
Connected governed pipeline	All stages complete and release is authorized	Accepted v0.3.8 candidate	Validated technical result
Planning artifacts	PRD, Epic/Feature/User Story, sizing, and evidence	Complete artifact set	Validated technical result
Delivery structure	Dynamic hierarchy with linked criteria	3 epics · 7 features · 11 stories · 11 criteria	Validated technical result
Citation governance	100% of run-specific citations receive an outcome	145/145	Validated technical result
Bidirectional integrity	Zero orphaned governed records	Zero orphans	Validated technical result
Hallucination control	Stage result ≤ 0.15	Accepted thresholds met	Validated technical result

Measure	Target	Current position	Classification
Time to first PRD draft	≥40% reduction from manual baseline	Baseline measurement required	Pilot target
Reviewer acceptance	≥80% without major rework	Pilot required	Pilot target
PM/TPM and engineering usefulness	≥4.0/5.0	Pilot survey required	Pilot target

Stakeholders

- Executive sponsor: business scope, resources, and promotion.
- AI PM/TPM: roadmap, governance, evidence, and delivery.
- Product managers: requirement and PRD approval.
- TPMs: dependencies, risks, and operating readiness.
- Engineering: feasibility and final estimation.
- Design/UX: personas, workflows, and experience questions.
- Business stakeholders: ambiguity and conflict decisions.
- Evaluation/governance: thresholds, evidence, and release recommendation.
- Security/privacy: data, credentials, access, and retention approval.

Risks and controls

Risk	Impact	Mitigation	Owner
Unsupported requirements	Teams plan unapproved work	Evidence-only prompts, citations, approval, reconciliation, and LLM evaluation	Program + evaluation owners
Incomplete or conflicting sources	False confidence or rework	Gap analysis, clarification gate, preserved viewpoints, fail-closed approval	Product owner
Automation bias	Generated content accepted without review	Human checkpoints, visible evidence, advisory sizing, accountable decisions	Program owner
Model or evaluator variability	Inconsistent release outcomes	Versioned prompts, repeated canaries, thresholds, regression tests, rollback	Evaluation owner
Evaluation cost growth	Unsustainable operating cost	Code controls first, targeted LLM evaluators, token/cache/cost monitoring	Operations owner
Sensitive-source exposure	Confidentiality or compliance failure	Approved storage, least privilege, controlled credentials, security review	Security owner
Workflow regression	Broken outputs or traceability	Candidate isolation, JSON exports, Git history, canaries, protected fallback	Workflow owner

Charter recommendation: Proceed to a limited, manually triggered pilot after submission evidence, ownership, security review, and manual baselines are complete. Preserve v0.3.7 as the protected fallback while v0.3.8 remains the candidate baseline.